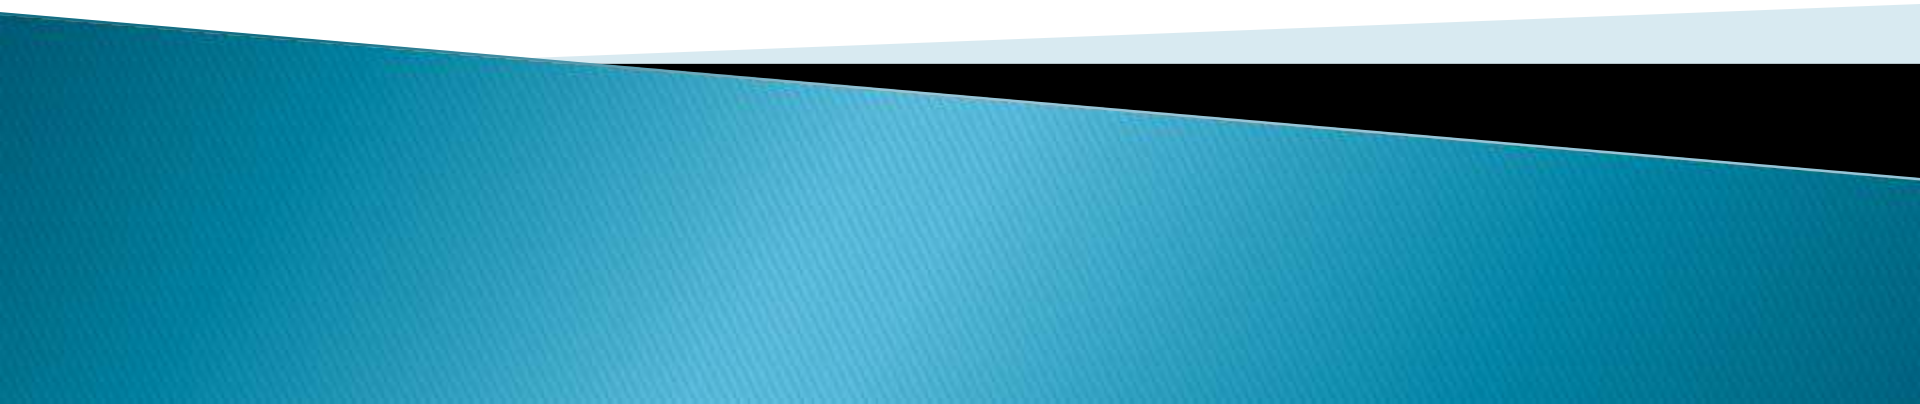



Processor Management

Damian Gordon



Processor Management

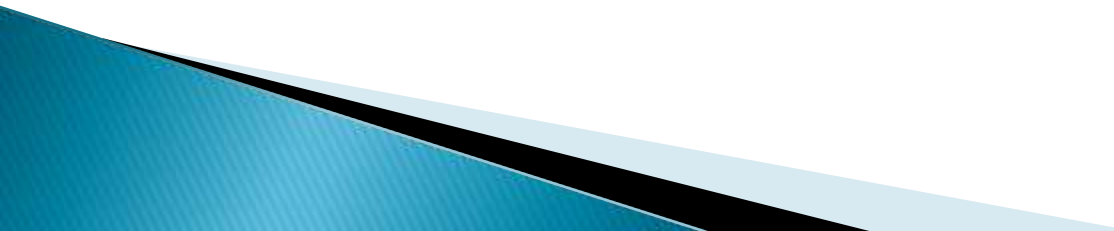
- ▶ If you were assembling some IKEA furniture, and following the step-by-step guide of 20 steps.
- ▶ Imagine you did step 1, then step 2, then step 3, suddenly the doorbell rang, and you are interrupted.



Processor Management

- ▶ You would stop what you are doing and deal with whoever is at the door.
- ▶ Then you would return to the assembly, check what step you were on, and then continue on from there (step 4).

Processor Management

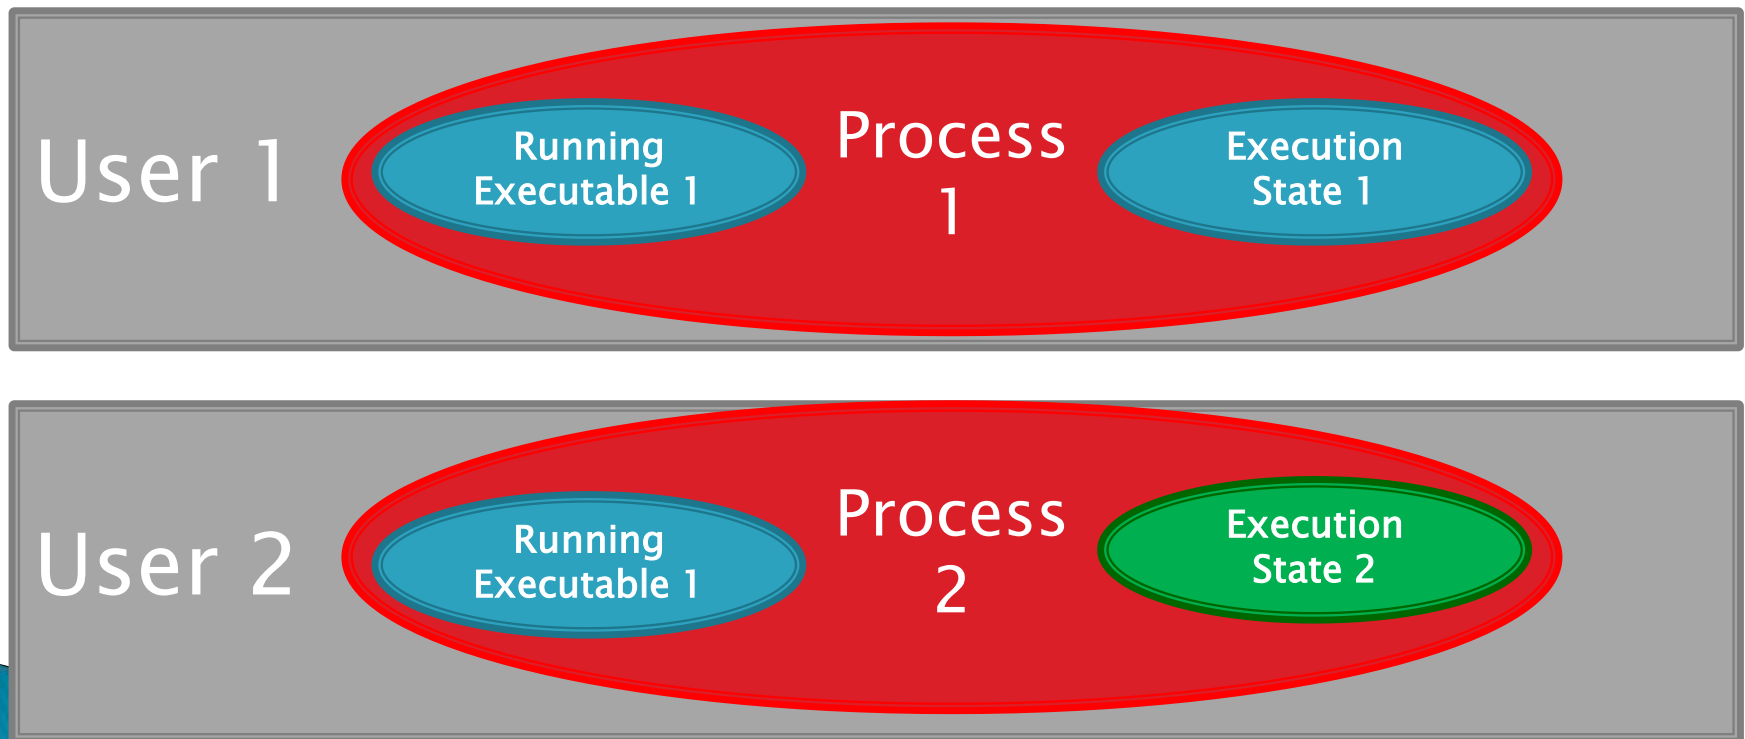
- ▶ This is how the computer runs multiple processes simultaneously ... you do a few steps on one process, get interrupted, work on other processes, then return to the first process, remembering where you left off from, and continue on.
 - ▶ This swapping of processes is called a **pre-emptive scheduling policy**.
- 

Processor Management

- ▶ A “Process” is the basic unit of execution in the operating system
- ▶ A “Process” is the name we give to a program when it is running in memory
 - So a “Program” is the complied executable
 - And a “Process” is the running executable with the execution state.


Processor Management

- ▶ The same program being run by two different users in the operating system are two different processes

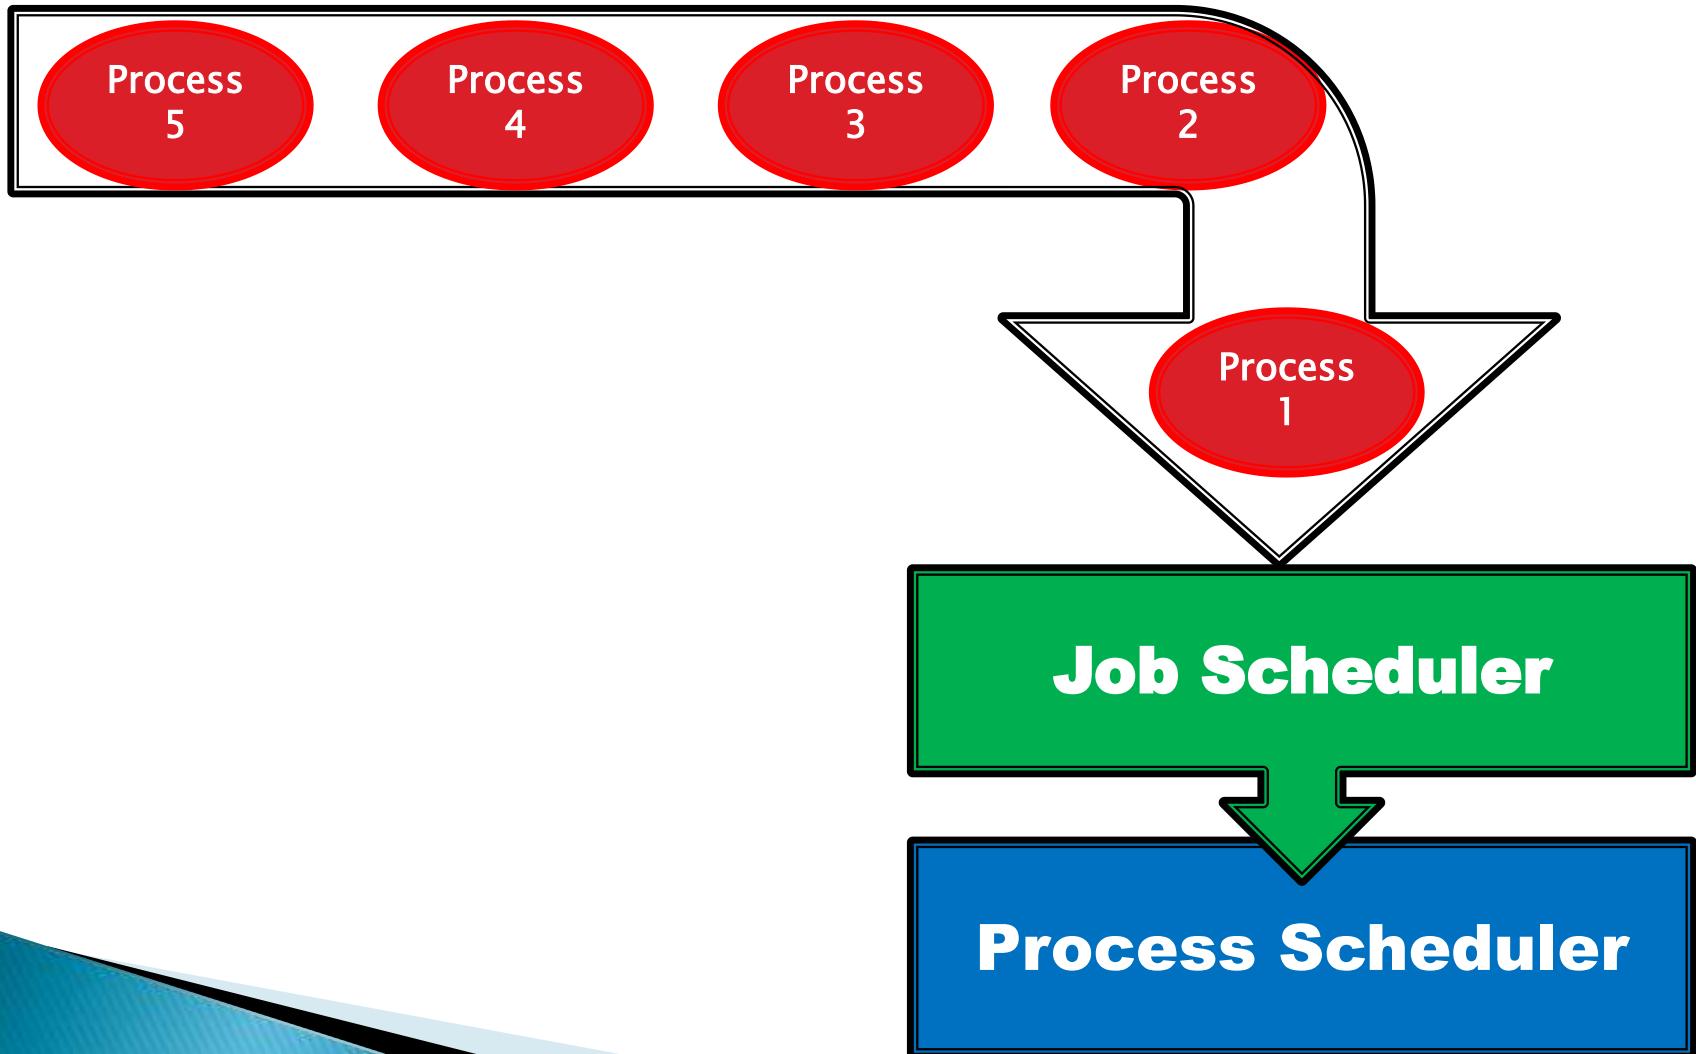


Processor Management

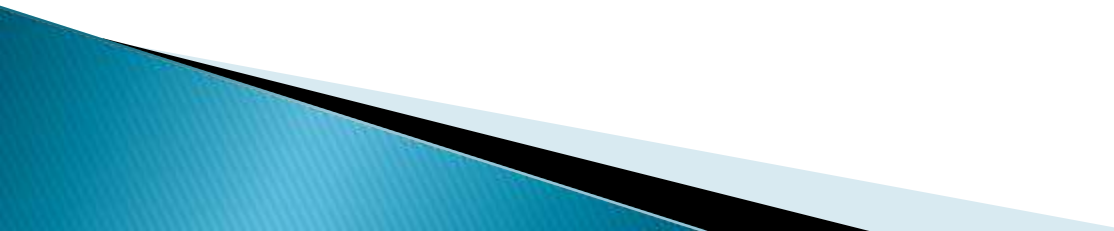
- ▶ The Processor Manager is made up of two sub-managers:



Processor Management



Job Scheduler

- ▶ We describe a collection of processes as a “job”.
 - ▶ When the operating system is presented with a queue of jobs or processes to run, the **Job Scheduler** has the task of deciding which order to run the jobs in.
 - ▶ The **Job Scheduler** wants to ensure that all components of the operating system are busy, and there is no component that is idle.
- 

Job Scheduler

▶ Let's think about this program:

```
PROGRAM PrintValue:
```

```
BEGIN
```

```
    Input A;
```

```
    Input B;
```

```
    C = A + B;
```

```
    D = A - B;
```

```
    Print "The sum of inputs is: ", C;
```

```
    Print "The Difference of inputs is: ", D;
```

```
END.
```



Job Scheduler

▶ Let's think about this program:

```
PROGRAM PrintValue:
```

```
BEGIN
```

```
Input A;
```

```
Input B;
```

I/O

```
C = A + B;
```

```
D = A - B;
```

```
Print "The sum of inputs is: ", C;
```

```
Print "The Difference of inputs is: ", D;
```

```
END.
```

Job Scheduler

- ▶ Let's think about this program:

```
PROGRAM PrintValue:
```

```
BEGIN
```

```
Input A;
```

```
Input B;
```

```
C = A + B;
```

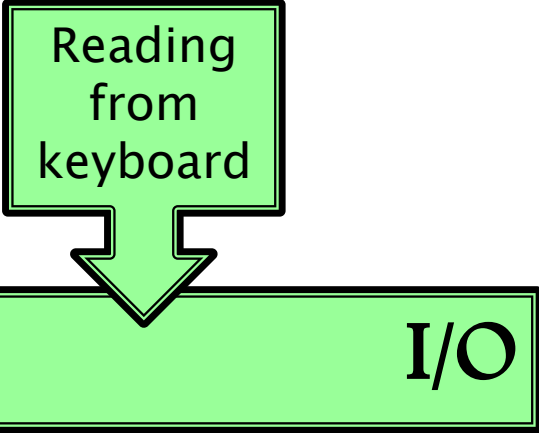
```
D = A - B;
```

```
Print "The sum of inputs is: ", C;
```

```
Print "The Difference of inputs is: ", D;
```

```
END.
```

Reading
from
keyboard

A green box with a black border contains the text "Reading from keyboard". A large green arrow points downwards from this box to a long green box below it. The long green box contains the code lines "Input A;" and "Input B;". To the right of the long green box, the text "I/O" is written in a large, bold, black font.

I/O

Job Scheduler

- ▶ Let's think about this program:

```
PROGRAM PrintValue:
```

```
BEGIN
```

```
Input A;
```

```
Input B;
```

I/O

```
C = A + B;
```

```
D = A - B;
```

```
Print "The sum of inputs is: ", C;
```

```
Print "The Difference of inputs is: ", D;
```

I/O

```
END.
```

Job Scheduler

- ▶ Let's think about this program:

```
PROGRAM PrintValue:
```

```
BEGIN
```

```
Input A;
```

```
Input B;
```

I/O

```
C = A + B;
```

```
D = A - B;
```

```
Print "The sum of inputs is: ", C;
```

```
Print "The Difference of inputs is: ", D;
```

I/O

```
END.
```



Writing
To
screen

Job Scheduler

- ▶ Let's think about this program:

```
PROGRAM PrintValue:
```

```
BEGIN
```

```
Input A;
```

```
Input B;
```

I/O

```
C = A + B;
```

```
D = A - B;
```

CPU

```
Print "The sum of inputs is: ", C;
```

```
Print "The Difference of inputs is: ", D;
```

I/O

```
END.
```

Job Scheduler

- ▶ Let's think about this program:

```
PROGRAM PrintValue:
```

```
BEGIN
```

```
Input A;
```

```
Input B;
```

```
C = A + B;
```

```
D = A - B;
```

```
Print "The sum of inputs is: ", C;
```

```
Print "The Difference of inputs is: ", D;
```

```
END.
```

I/O


Doing
computation

CPU

I/O


Job Scheduler

- ▶ Some programs do a lot of I/O, e.g. Graphics programs:




Job Scheduler

- ▶ Others do a lot of computation, but little I/O, e.g. maths programs:




Job Scheduler

- ▶ If the job scheduler gets jobs like this:




Job Scheduler

- ▶ If the job scheduler gets jobs like this:




Job Scheduler

- ▶ So the job scheduler will take these jobs:

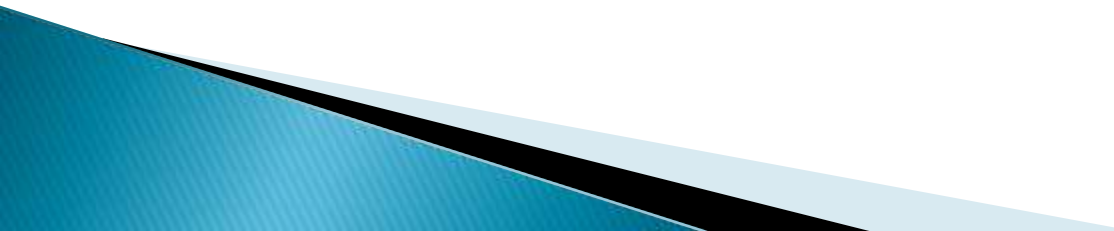


Job Scheduler

- ▶ And swap them around, and pass them onto the Process Scheduler



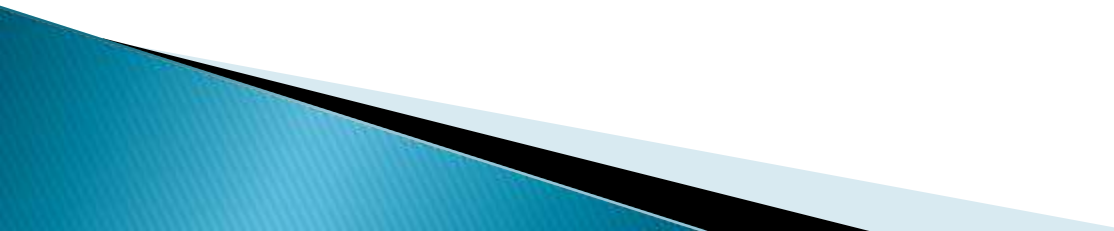
Job Scheduler

- ▶ So in this case the Job Scheduler wants to balance the jobs coming in, so that the components of the operating system are all kept busy.
 - ▶ If we don't change the order of job, the CPU will be very busy, but the I/O component will be idle, and then vice versa.
 - ▶ This is not an optimal use of resources, so we swap them around.
- 


Job Scheduler

- ▶ We'll call jobs that are CPU-bound **Batch Jobs**, and we'll call jobs that have a lot of I/O operations **Interactive Jobs**.

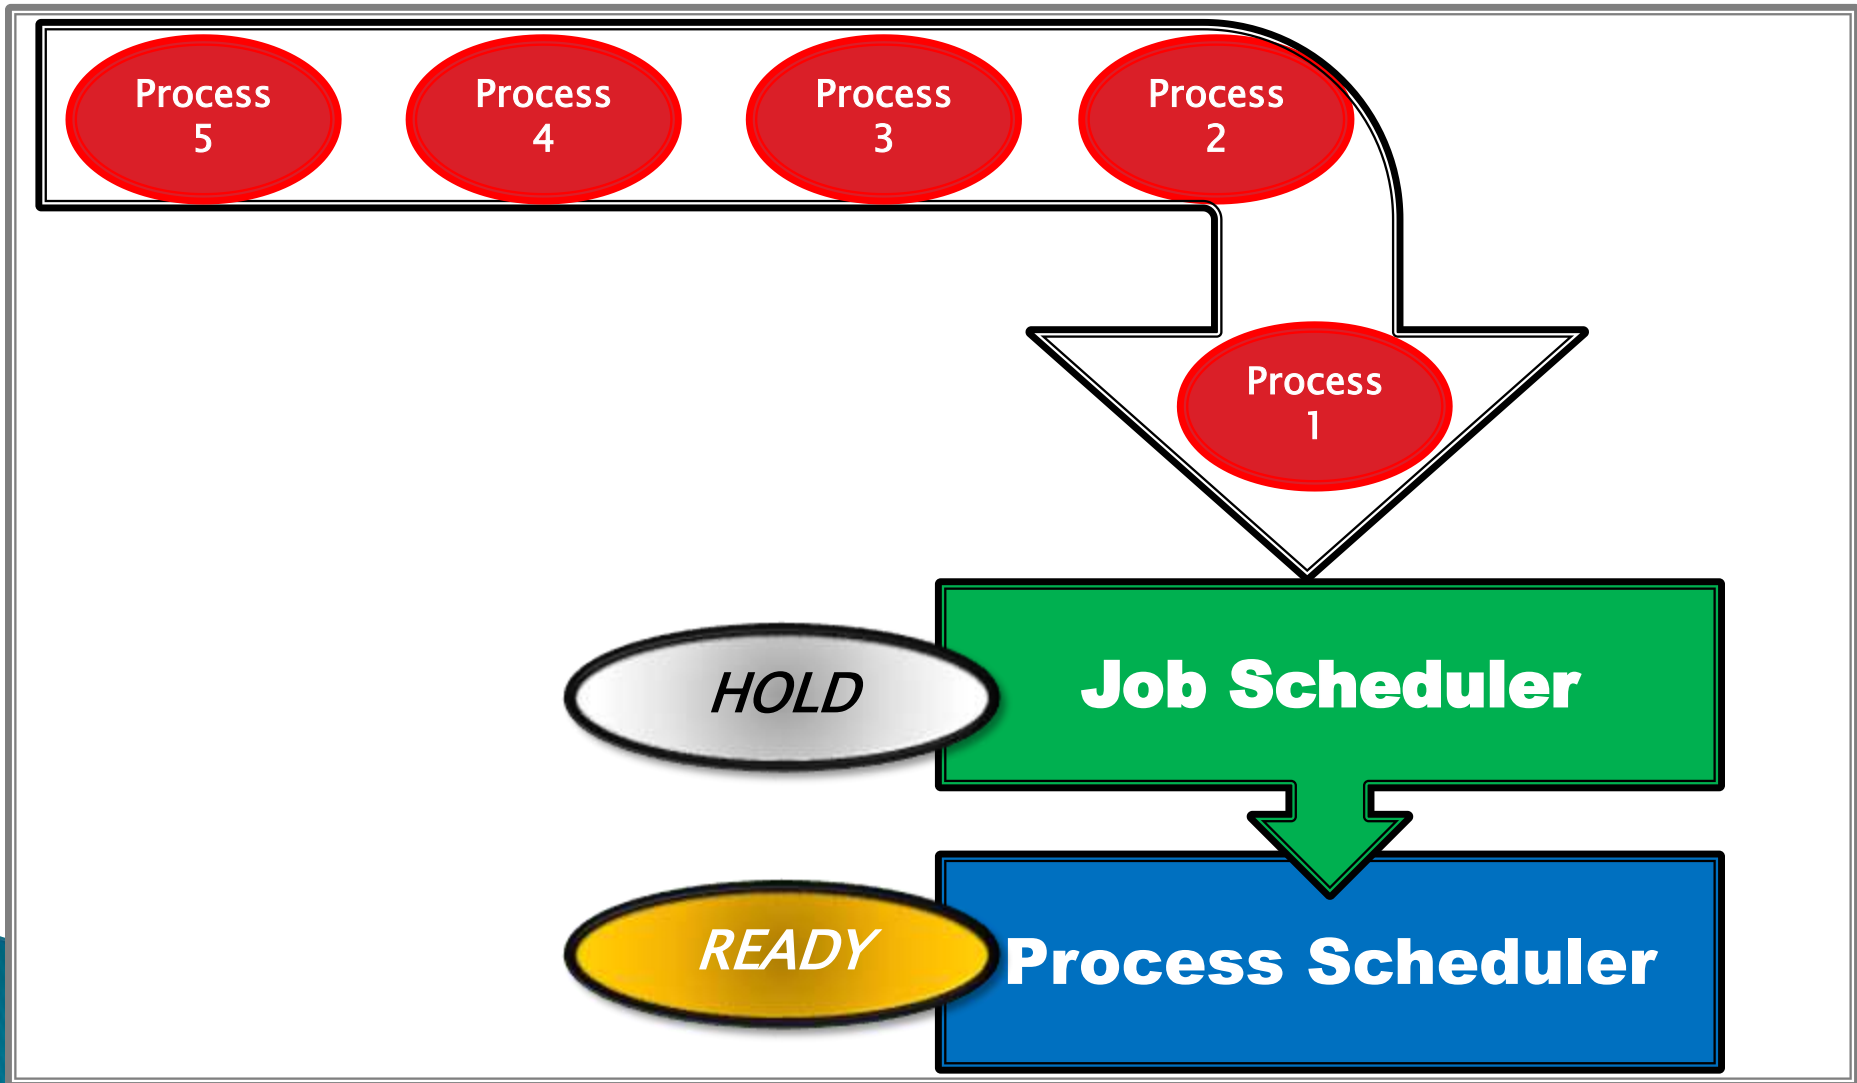
Job Scheduler

- ▶ This is the first level of swapping of processes that will occur, and more will be done by the **Process Scheduler**. The **Job Scheduler** is looking at jobs (or groups of processes), and looking at the whole process from a high-level.
 - ▶ For obvious reasons, the Job Scheduler is also called the **High-Level Scheduler**.
- 

Job Scheduler

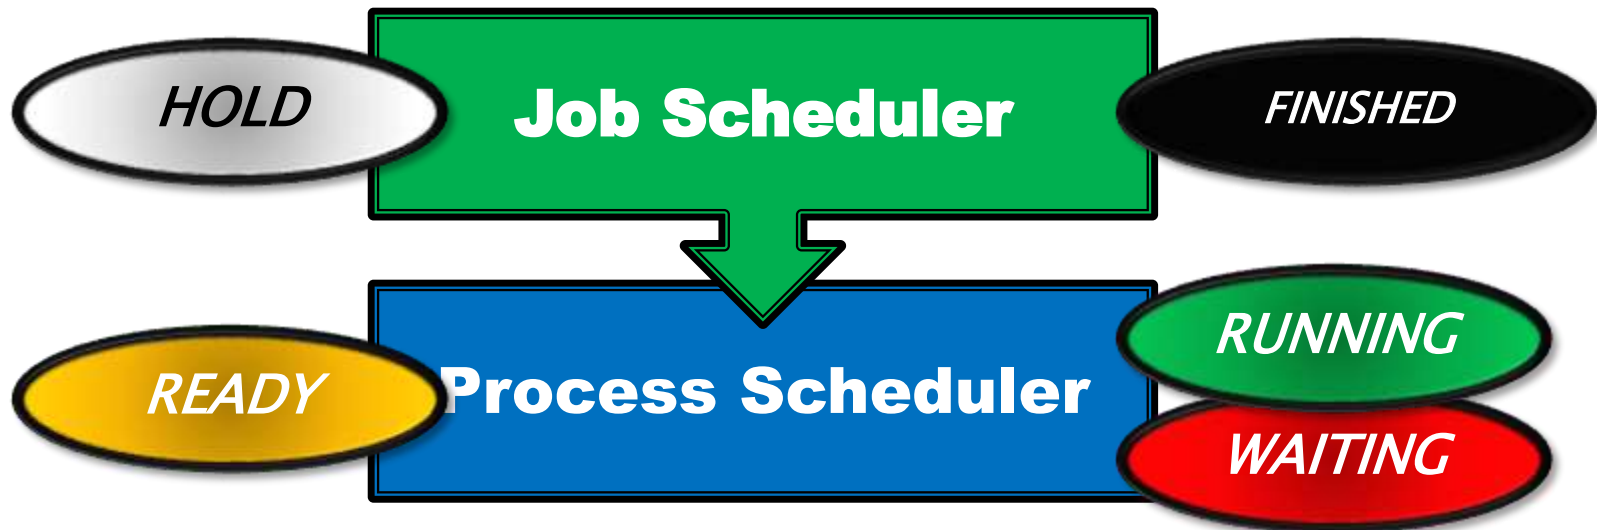
- ▶ Every process has a field that records their current status, called the **Process Status**.
 - ▶ When the process is first passed to the **Job Scheduler** from the operating system, its status is always set as **HOLD**.
 - ▶ When the **Job Scheduler** passes the process onto the **Process Scheduler**, its status is always changed to **READY**.
- 

Processor Management

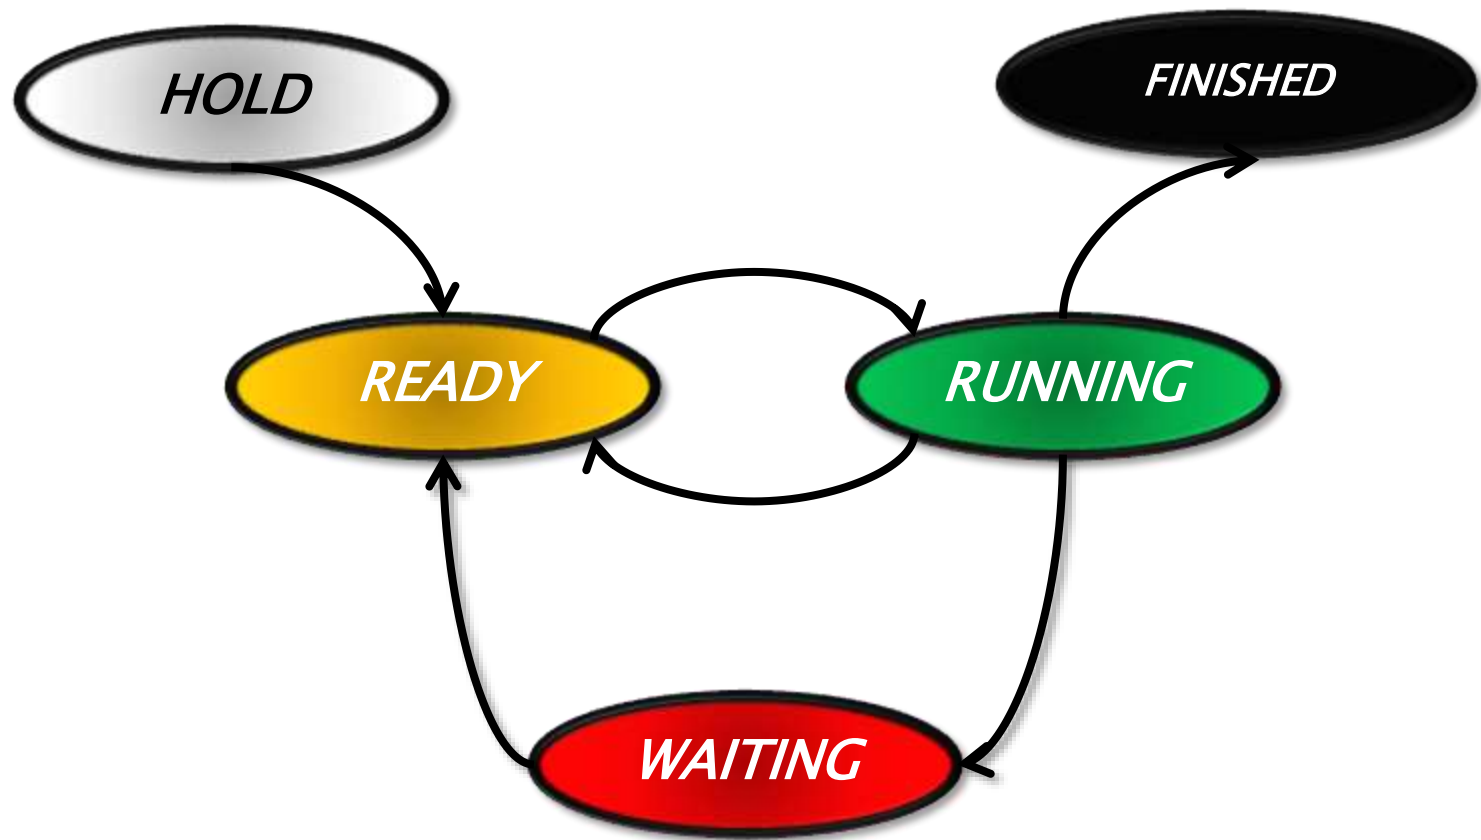


Processor Management

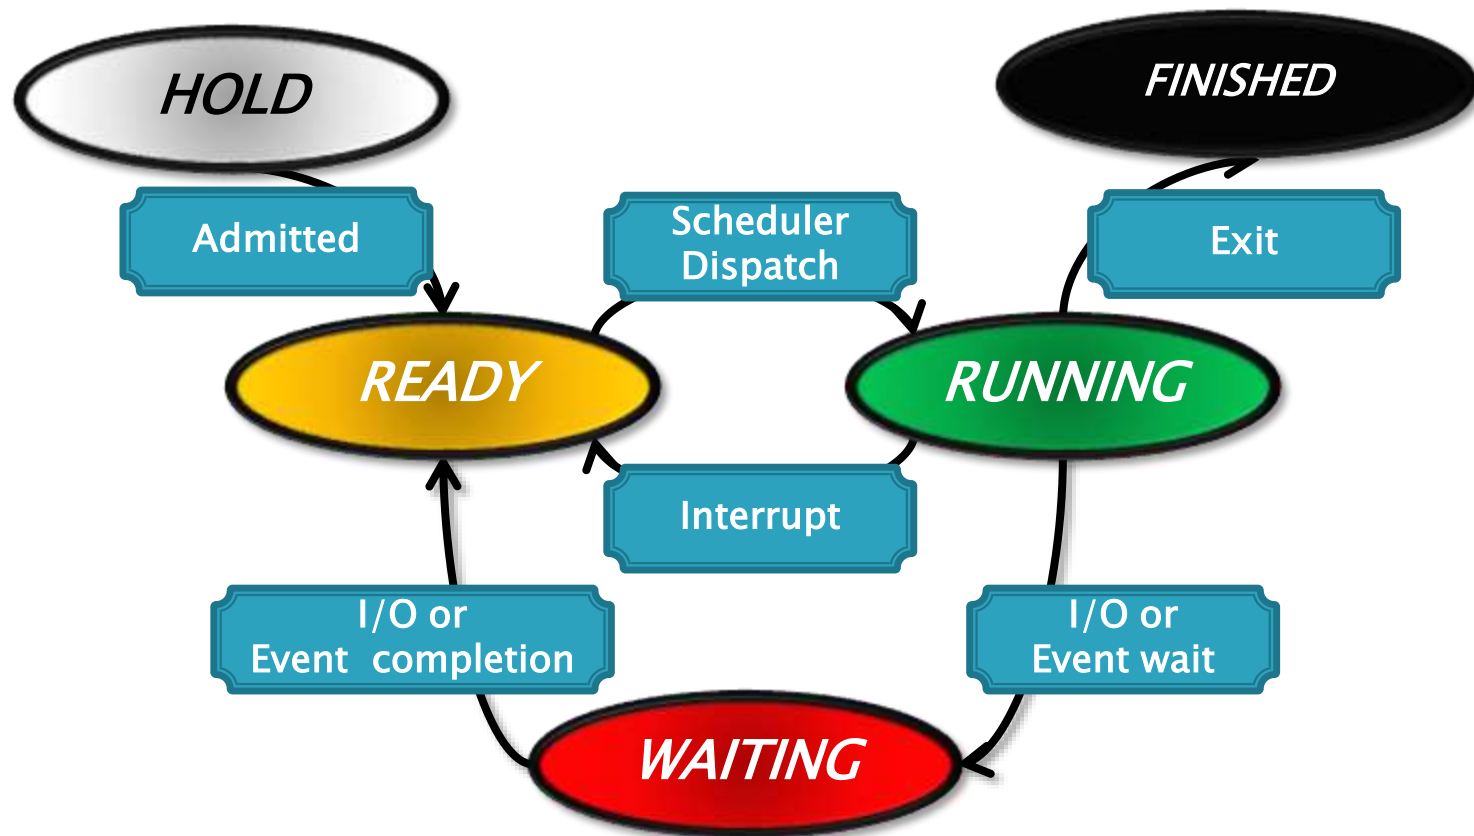
- ▶ Other statuses that a process can have are:



Process Scheduler



Process Scheduler




Process Scheduler

- ▶ Think about traffic lights:

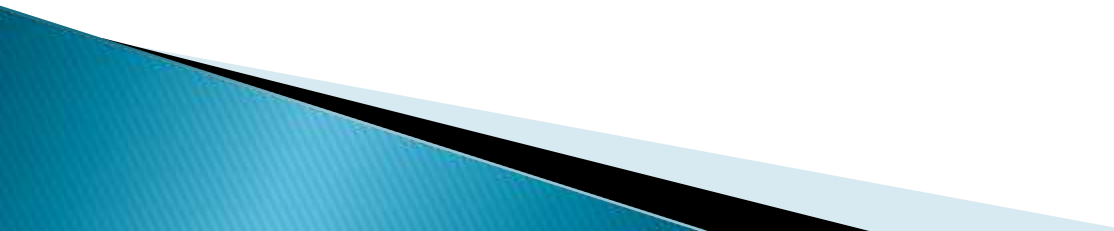


Process Scheduler

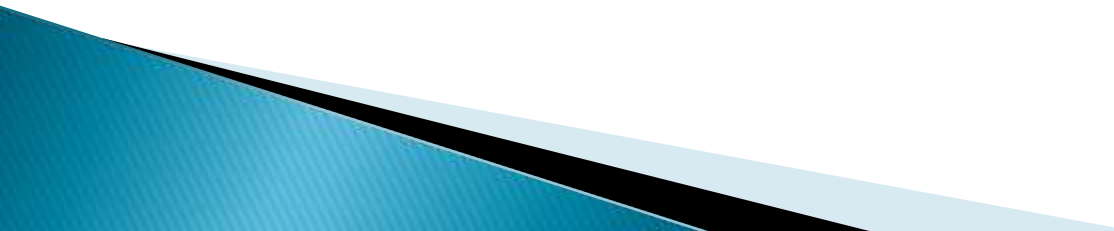
- ▶ This is the sequence:



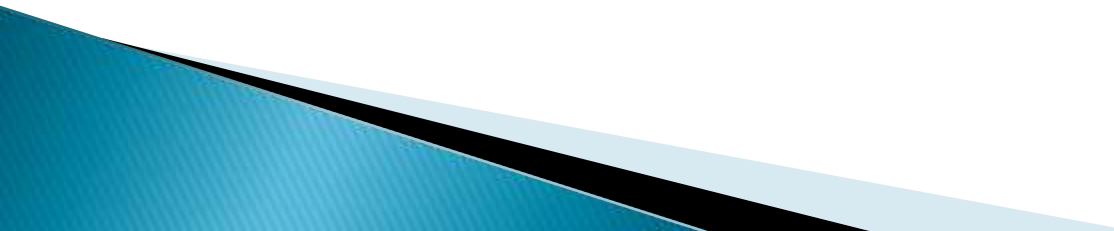
Process Scheduler

- ▶ As mentioned previously the process is first passed to the **Job Scheduler** from the operating system, its status is always set as **HOLD**. When the **Job Scheduler** passes the process onto the **Process Scheduler**, its status is always changed to **READY**.
- 

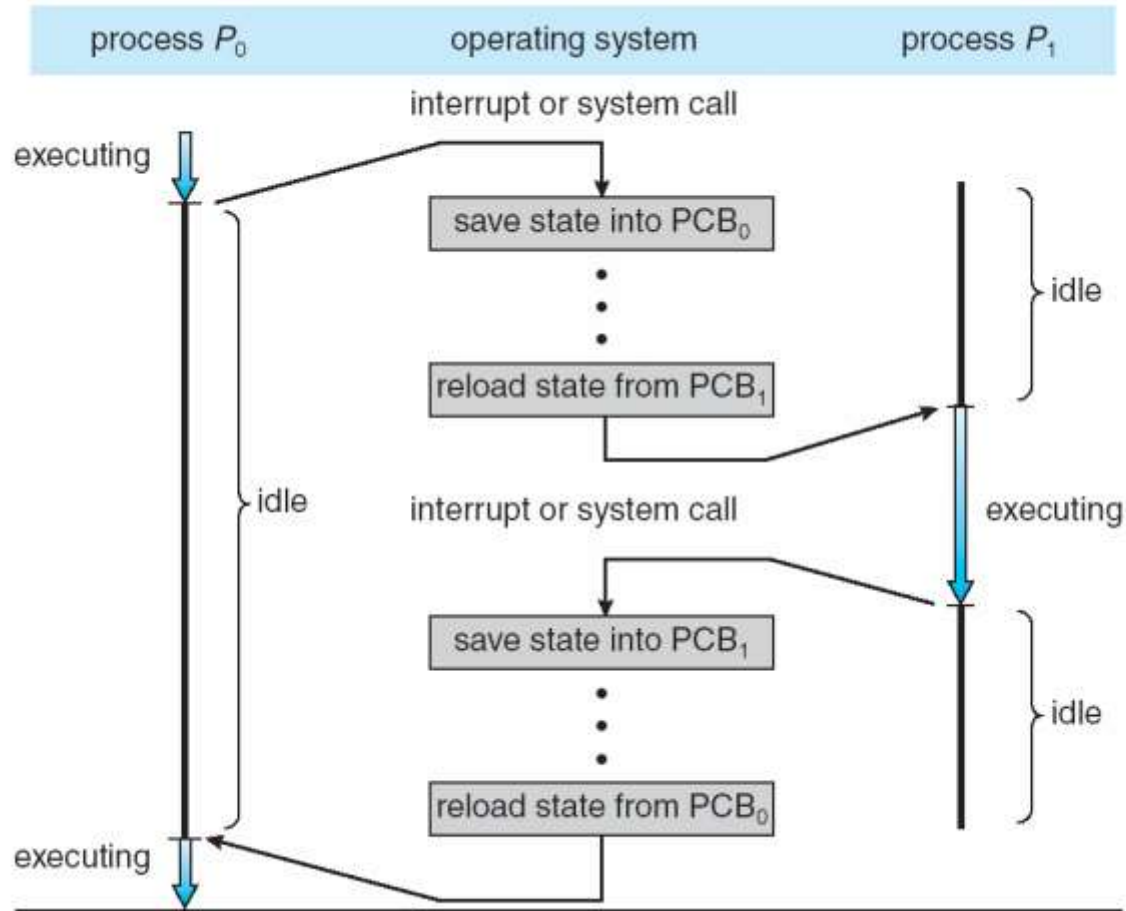
Process Scheduler

- ▶ When the CPU is available, the **Process Scheduler** will look at all of the upcoming processes and will select one of them (based on a predefined algorithm) and assuming there is memory free, the process will start running and its status is changed to **RUNNING** by the **Process Scheduler**.
- 

Process Scheduler

- ▶ After a predefined time, the process will be interrupted, and another process will take over the CPU, and the interrupted process will have its status changed to **READY** by the **Process Scheduler**. We will remember that this swapping of processes is called a **pre-emptive scheduling policy**.
 - ▶ When is CPU is available for this process again the process status be changed to **RUNNING** by the **Process Scheduler**.
- 

Process Scheduler



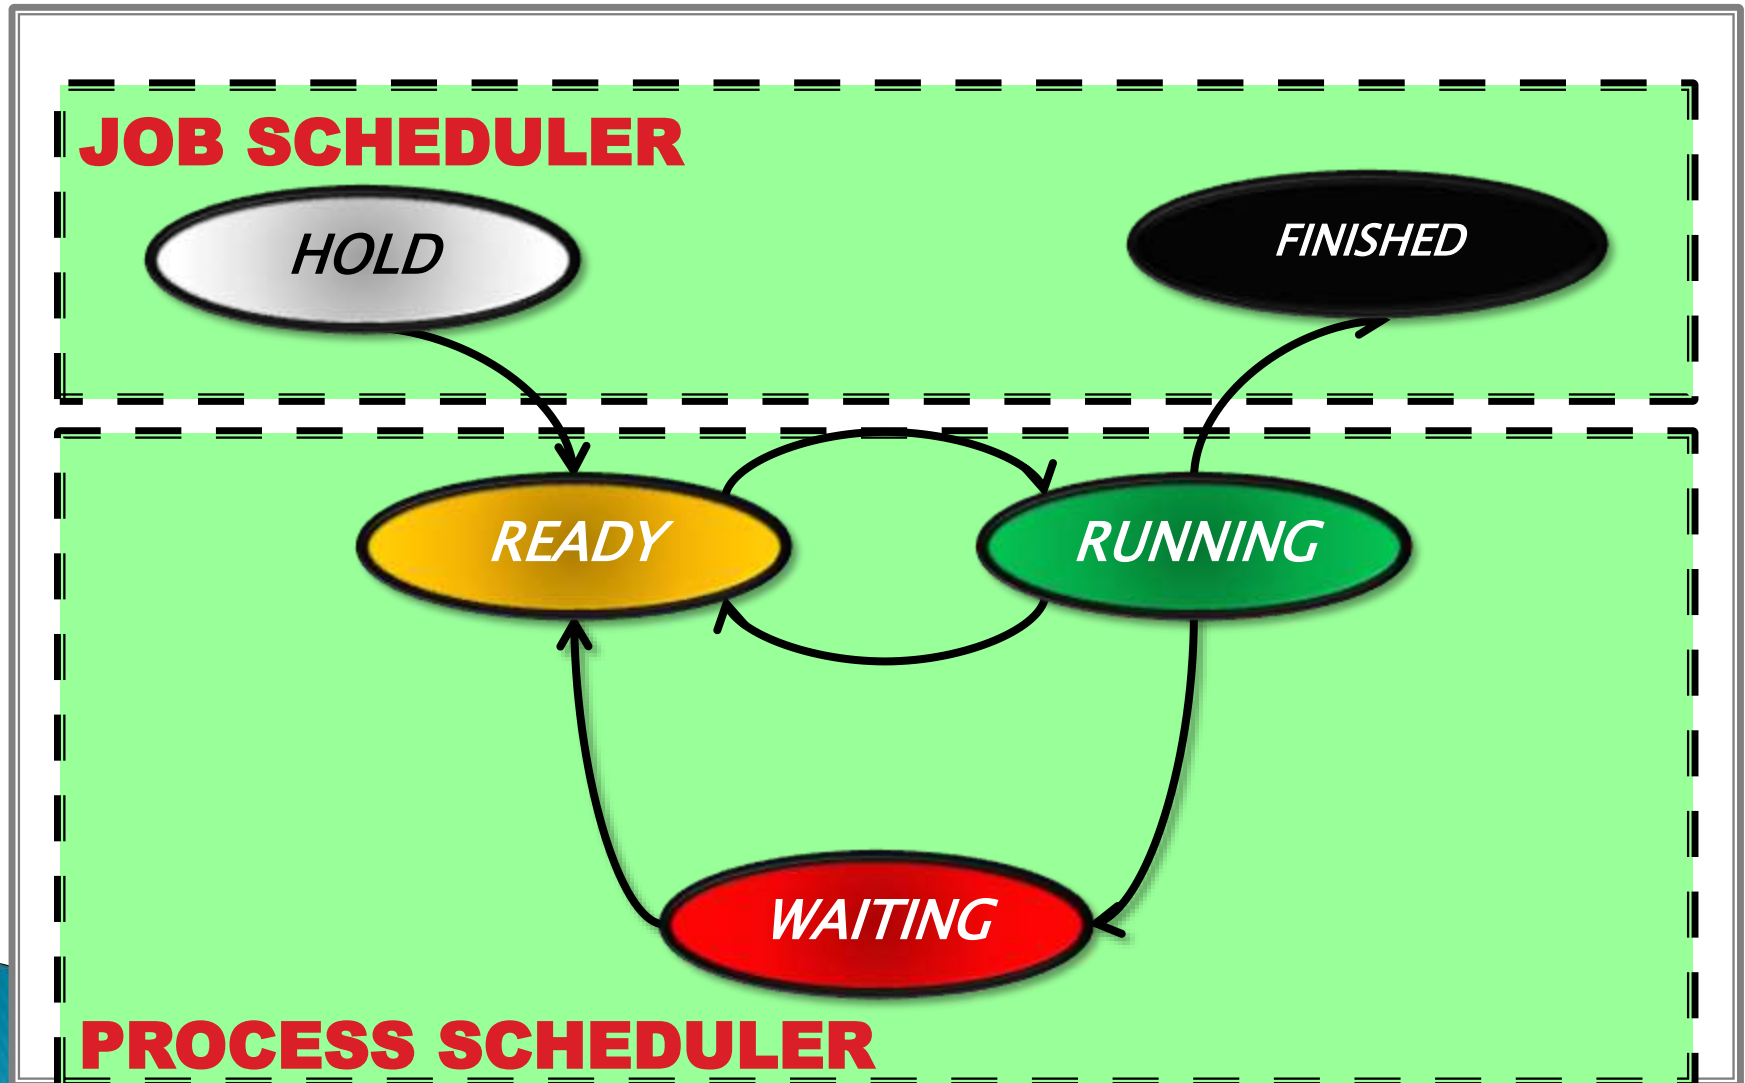
Process Scheduler

- ▶ If the process has to wait for some I/O from the user or some other event, it is put in a status of **WAITING** by the **Process Scheduler**.
- ▶ When the I/O device lets the **Process Scheduler** know that the I/O is completed, the process status will be changed to **READY** by the **Process Scheduler**.

Process Scheduler

- ▶ Finally the process status will be changed to **FINISHED** either when the process has successfully completed, or if an error occurs and the process has to terminate prematurely. The status change is usually handled by the **Process Scheduler** informing the **Job Scheduler** of the process completion, and the **Job Scheduler** of changing the status.

Processor Management



Process Control Block (PCB)

- ▶ The Process status is only one of a collection of descriptors that are associated with a process.
- ▶ Each process has a data structure called a Process Control Block (PCB) associated with it, rather like a passport.

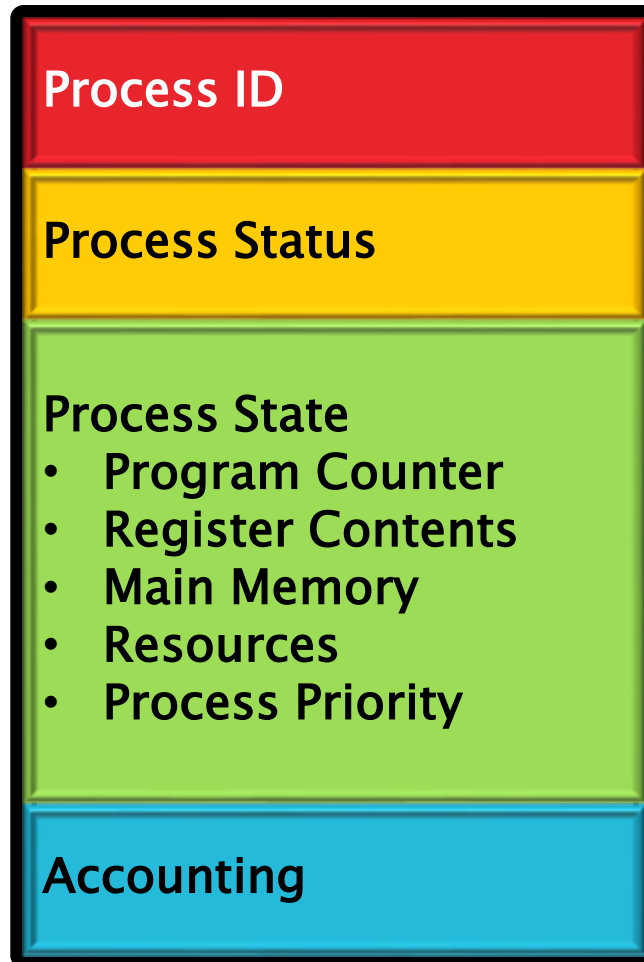
Process Control Block (PCB)

Operating System



Process Control Block
(PCB)

Process Control Block (PCB)



Process Control Block (PCB)

▶ **PROCESS IDENTIFIER**

- Each process is uniquely identified by both the user's identification, and a pointer connecting it to its descriptor.

Process Control Block (PCB)

▶ PROCESS STATUS


- The current status of the process, it is either:



Process Control Block (PCB)

▶ PROCESS STATE



- Program counter
- Record the current value of the program counter



Process Control Block (PCB)

▶ PROCESS STATE


- Register Contents
- Record the values of the data registers



Process Control Block (PCB)

▶ PROCESS STATE

- Main Memory
- Record all important information from memory, including most importantly the process location.



Process Control Block (PCB)

▶ PROCESS STATE


- Resources
- Record the resources that have been allocated to this job, including files, disk drives, and printers.

ID	TYPE	DETAILS
1	FILE	TXT file starting at memory address 0x456
2	FILE	DAT file starting at memory address 0x087
3	DISK	Disk Drive 4
4	FILE	TXT file starting at memory address 0x673
5	PRINTER	Printer at IP address 172.242.34.65

Process Control Block (PCB)

▶ PROCESS STATE


- Process Priority
- The process is assigned a priority, and if the operating system using priorities to schedule processes.



Process Control Block (PCB)

▶ PROCESS STATE


- Process Priority
- The process is assigned a priority, and if the operating system uses priorities to schedule processes.



Process Control Block (PCB)

▶ PROCESS STATE

- Process Priority
- The process is assigned a priority, and if the operating system uses priorities to schedule processes.

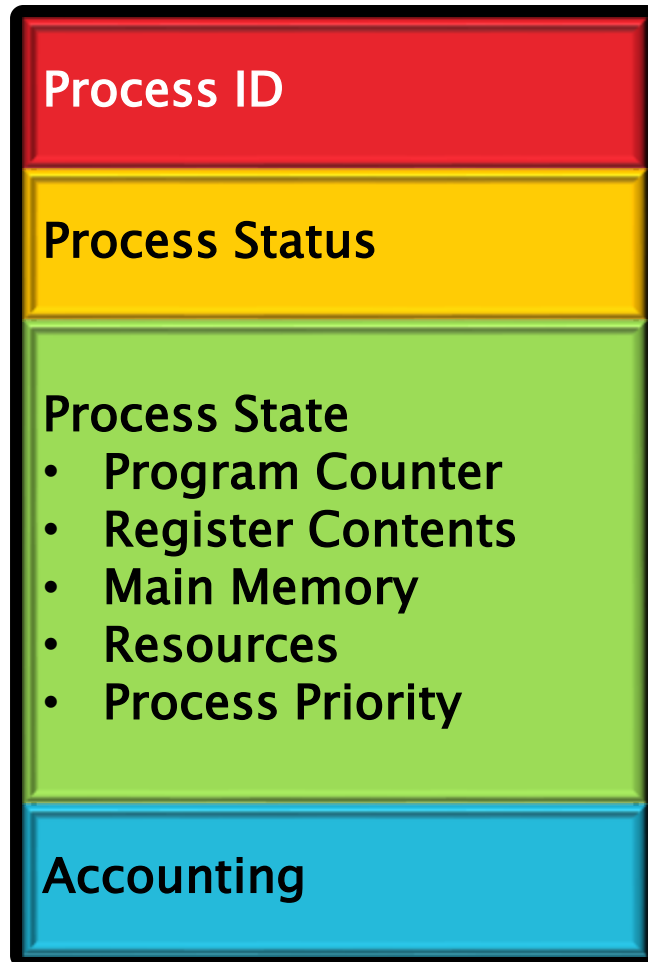


Process Control Block (PCB)

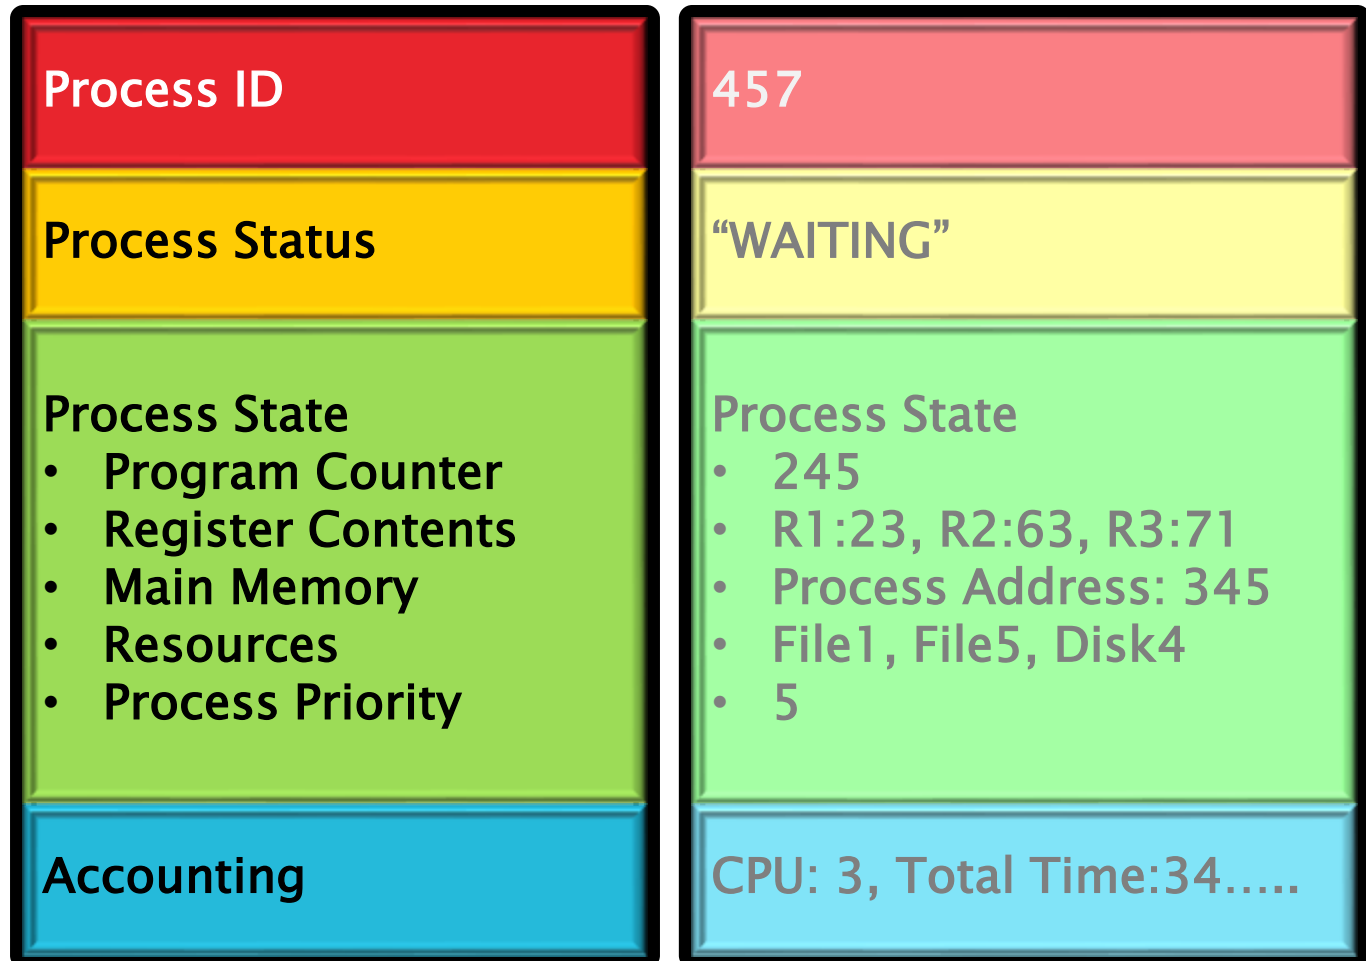
▶ ACCOUNTING

- This section records some of the performance statistics associated with this process, including:
 - CPU time used so far
 - Time process has been in the system
 - Time taken in memory (Main and secondary)
 - Space taken in memory (Main and secondary)
 - System programs used, compilers, editors, etc.
 - Number and type of I/O operations
 - Time spent waiting for I/O operations completion
 - Number of Input records read and Output records written

Process Control Block (PCB)



Process Control Block (PCB)



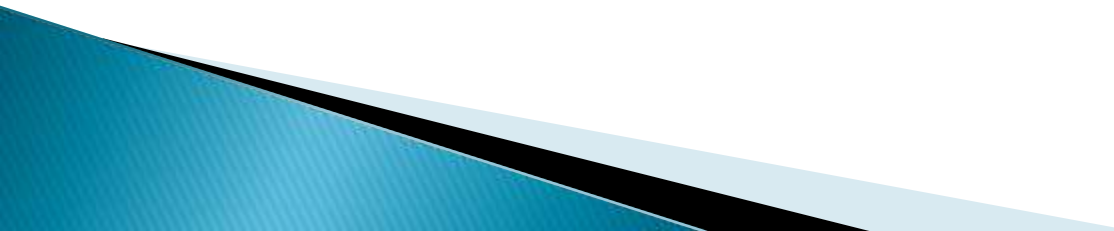
Process Scheduling Policies



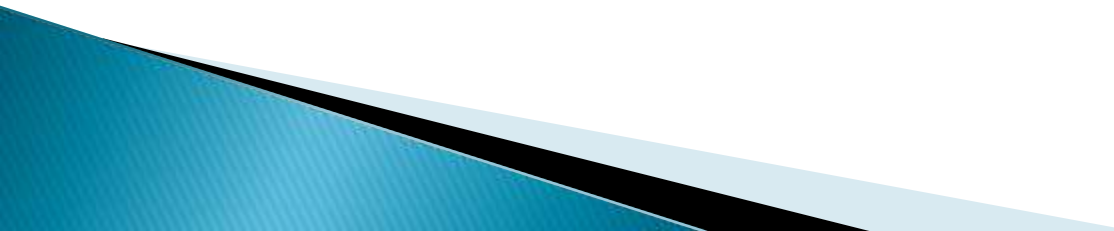
Process Scheduling Policies

- ▶ What are good policies to schedule processes?

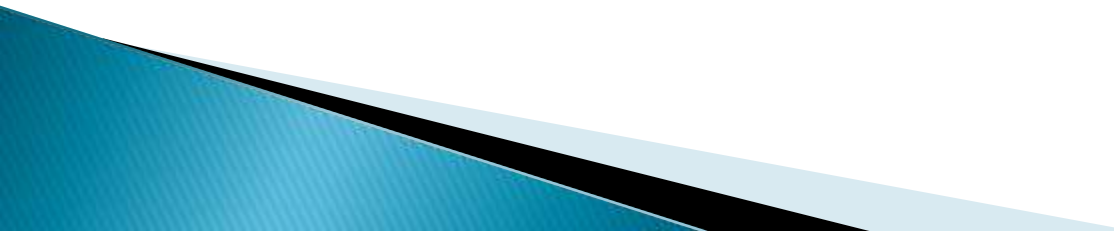
Process Scheduling Policies

- ▶ **MAXIMUM THROUGHPUT**
 - ▶ Get as many jobs done as quickly as possible.
 - ▶ There are several ways to achieve this, e.g. run only short jobs, run jobs without interruptions.
- 

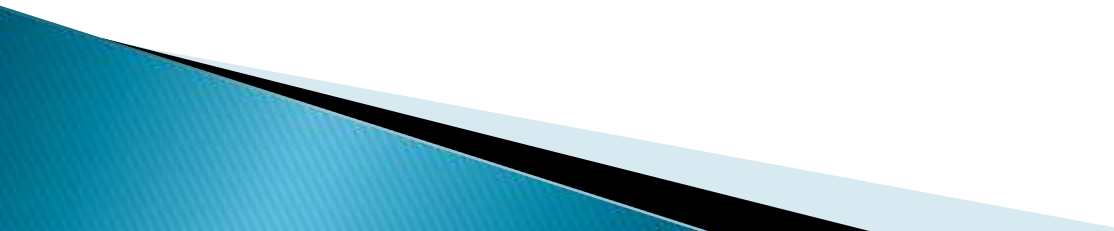
Process Scheduling Policies

- ▶ **MINIMIZE RESPONSE TIME**
 - ▶ Ensure that interactive requests are dealt with as quickly as possible.
 - ▶ This could be achieved by scheduling just with a lot of I/O jobs first, and leave the computational jobs for later.
- 

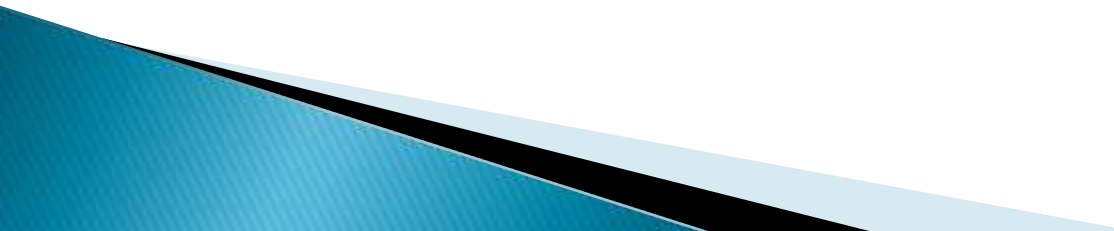
Process Scheduling Policies

- ▶ **MINIMIZE TURNAROUND TIME**
 - ▶ Ensure that jobs are completed as quickly as possible.
 - ▶ This could be achieved by scheduling just with a lot of computation jobs first, and leave the I/O jobs for later, so there is no user delays.
- 

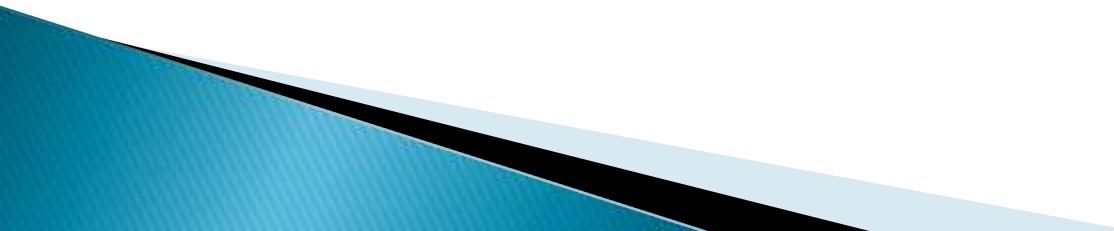
Process Scheduling Policies

- ▶ **MINIMIZE WAITING TIME**
 - ▶ Move jobs out of the READY status as soon as possible.
 - ▶ This could be achieved by reduce the number of users allowed on the system, so that the CPU would be available whenever a job enters the READY status.
- 

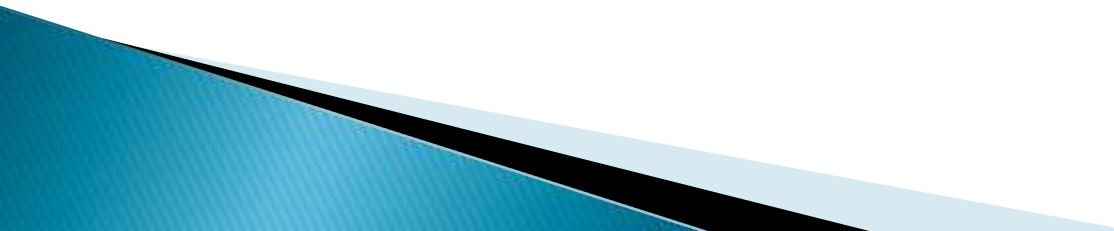
Process Scheduling Policies

- ▶ **MINIMIZE WAITING TIME**
 - ▶ Move jobs out of the READY status as soon as possible.
 - ▶ This could be achieved by reduce the number of users allowed on the system, so that the CPU would be available whenever a job enters the READY status.
- 

Process Scheduling Policies

- ▶ **MAXIMISE CPU EFFICIENCY**
 - ▶ Keep the CPU busy 100% of the time.
 - ▶ This could be achieved by scheduling just with a lot of computation jobs, and never run the I/O jobs.
- 

Process Scheduling Policies

- ▶ **ENSURE FAIRNESS FOR ALL JOBS**
 - ▶ Give everyone an equal amount of CPU time and I/O time.
 - ▶ This could be achieved by giving all jobs equal priority, regardless of it's characteristics.
- 

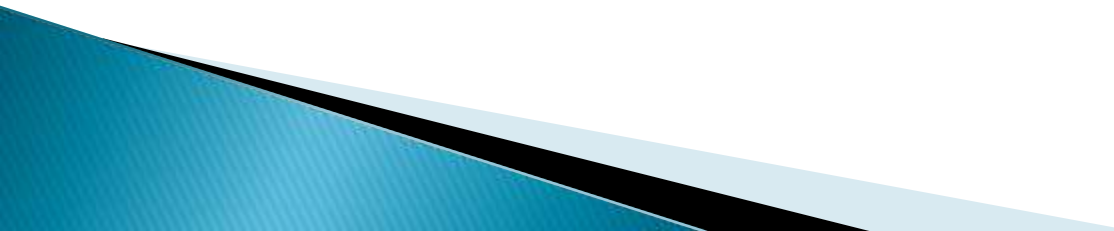
Process Scheduling Algorithms




Process Scheduling Algorithms

- ▶ Here are six commonly used process scheduling algorithms

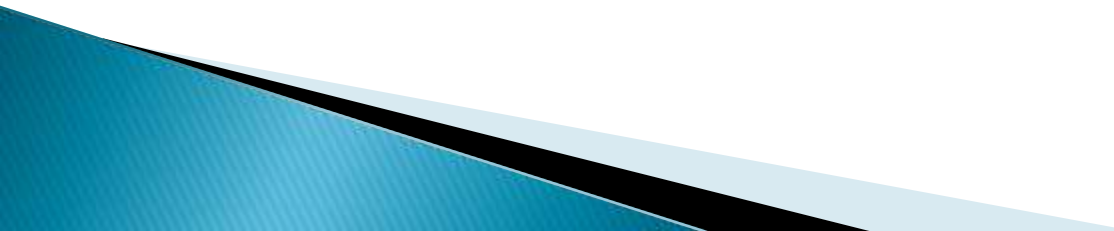
Process Scheduling Algorithms

- ▶ **First Come, First Served (FCFS)**
 - ▶ A very simple algorithm that uses a FIFO structure.
 - ▶ Implemented as a non-pre-emptive scheduling algorithm.
 - ▶ Works well for **Batch Processes**, where users don't expect any interaction.
- 

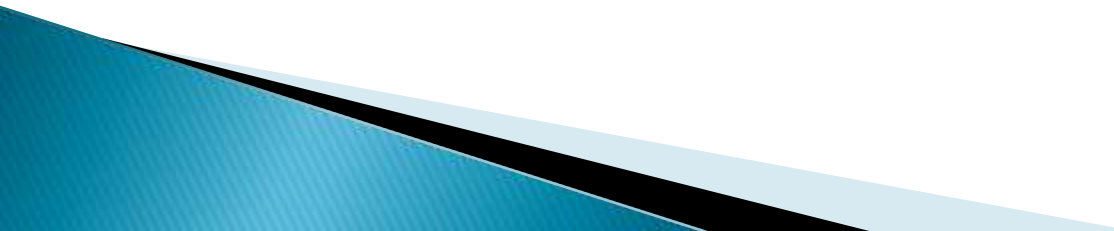
Process Scheduling Algorithms

- ▶ **Shortest Job Next (SJN)**
 - ▶ Also called Shortest Job First (SJF).
 - ▶ A very simple algorithm that schedules processes based on CPU cycle time.
 - ▶ Implemented as a non-pre-emptive scheduling algorithm.
 - ▶ Works well for **Batch Processes**, where it is easy to estimate CPU time.
- 

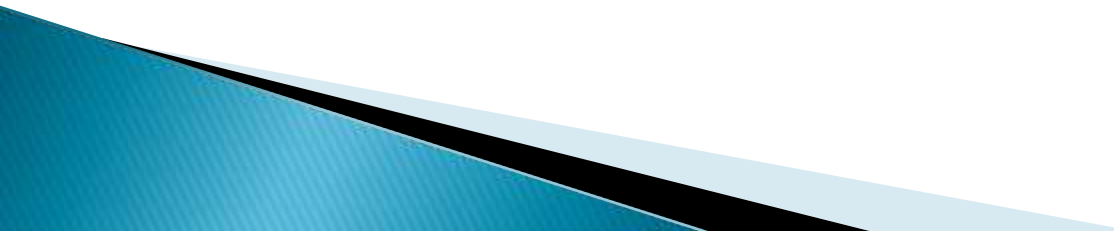
Process Scheduling Algorithms

- ▶ **Priority Scheduling**
 - ▶ An algorithm that schedules processes based on priority.
 - ▶ Implemented as a non-pre-emptive scheduling algorithm.
 - ▶ One of the most common algorithms used in systems that are mainly **Batch Processes**.
 - ▶ If two jobs come in of equal priority are READY, it works on a FIRST COME, FIRST SERVED basis.
- 

Process Scheduling Algorithms

- ▶ **Shortest Remaining Time (SRT)**
 - ▶ A pre-emptive scheduling version of the Shortest Job Next (SJN) algorithm.
 - ▶ An algorithm that schedules processes based on the one which is nearest to completion.
 - ▶ It can only be implemented on systems that are only **Batch Processes**, since you have to know the CPU time required to complete each job.
- 

Process Scheduling Algorithms

- ▶ **Round Robin**
 - ▶ A pre-emptive scheduling algorithm that is used extensively in interactive systems
 - ▶ All active processes are given a pre-determined slice of time ("*time quantum*").
 - ▶ Choosing the time quantum is the key decision, for interactive systems the quantum must be small, whereas for batch systems it can be longer.
- 

Process Scheduling Algorithms

- ▶ **Multi-Level Queues**
 - ▶ This isn't really a separate scheduling algorithm, it can be used with others.
 - ▶ Jobs are grouped together based on common characteristics.
 - ▶ For example, CPU-bound jobs based in one queue, and the I/O-bound jobs in another queue, and the process scheduler can select jobs from each queue based on balancing the load.
- 